

Algorand Consensus

Giuseppe Persiano

Università di Salerno
giuper@gmail.com

An overview

Algorand Consensus: the executive summary



Pure Proof of Stake (PPoS) Byzantine Agreement

- can tolerate malicious users
- no need for a central authority, as long as a **super-majority** of the **stake** is in non-malicious hands.
- extremely fast (currently, one new **block** every ≈ 4 sec)
- **fork** with probability $< 10^{-18}$
- immediate **finality**
- no delegation or binding of the **stake**
- requires minimal computational power per node
- carbon neutral

► Source: https://developer.algorand.org/docs/get-details/algorand_consensus/

Digital currency

Double spending

Each digital coin can only be spent once.

- **transaction based:** transactions take coins as input and output coins need to keep track of which coins have not been spent yet
coins input to a transaction are burned
 - ▶ Bitcoin keeps a list of *unspent transaction outputs* UTXO
- **account based:** transactions move coins from one account to another need to keep track of number of coins associated with each account

the concept of a *transaction* is basic in both models

Current implementations

- Physical coins and bills:

- ▶ each individual is an account
- ▶ an individual owns the coins in his pocket/wallet
- ▶ a transaction moves money from one pocket to another

- Bank accounts:

- ▶ a transaction moves money from one account to another
- ▶ the bank keeps track of the balance in each account
- ▶ **private, centralized and offline** ledger

We want **public, decentralized, on-line ledgers**

Consensus on who owns money

- Physical coins and bills:

- ▶ Consensus is trivial: we all agree that what's in my pocket is mine

- Bank accounts:

- ▶ Because the **BANK** says so

Decentralized OnLine Consensus

Decentralized OnLine Consensus

Consensus on transactions

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*
- A block is a set of *transactions*

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*
- A block is a set of *transactions*
- Nodes of the network must reach consensus on the next block

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*
- A block is a set of *transactions*
- Nodes of the network must reach consensus on the next block
- A new block is proposed by one node of the network

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*
- A block is a set of *transactions*
- Nodes of the network must reach consensus on the next block
- A new block is proposed by one node of the network
- Which node?

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*
- A block is a set of *transactions*
- Nodes of the network must reach consensus on the next block
- A new block is proposed by one node of the network
- Which node?
 - ▶ The selected node is the one that can provide the solution to a *public* puzzle

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*
- A block is a set of *transactions*
- Nodes of the network must reach consensus on the next block
- A new block is proposed by one node of the network
- Which node?
 - ▶ The selected node is the one that can provide the solution to a *public* puzzle

Decentralized OnLine Consensus

Consensus on transactions

- A blockchain is a *public* ledger organized as a sequence of *blocks*
- A block is a set of *transactions*
- Nodes of the network must reach consensus on the next block
- A new block is proposed by one node of the network
- Which node?
 - ▶ The selected node is the one that can provide the solution to a *public* puzzle

Puzzles

- *Efficiency*: the cost of solving a puzzle or realizing that we cannot solve the puzzle
 - ▶ green vs non-green blockchains
- *Multiple Solvers*: the puzzle can be solved concurrently by more than one node or each puzzle can only be solved by one node.
 - ▶ block finality
- The block proposed by the selected node *must* be verified by the other nodes

Distributed vs Centralized

- Centralized consensus is easy
 - ▶ Banks have been successfully operating for centuries
- Decentralized consensus is hard
 - ▶ No trusted party
 - ▶ Only trust that a (super-)majority is honest
 - ▶ Majority of what??? Sybil attack
 - ▶ Bitcoin: computing power
 - ▶ Algorand: tokens
 - ▶ Cannot be faked and does not need identity management

Bitcoin vs Algorand

Bitcoin's Puzzle

Proof of Work:

- hard-to-invert hash function
- probability of selection proportional to computing power
- the same puzzle might be solved by two (or more) nodes

Algorand's Puzzle

Proof of Stake:

- each token (*algo*) throws a die with prob $1/N$ of winning
- winning token can be shown
- probability of winning proportional to number of tokens held
- probability of two algos throwing the winning die is about 10^{-18}

Proof of Work

Dwork and Naor, 1992

Hashcash by Adam Back, 1997

Combatting junk mail

- easy to compute *hash function* $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$
- infeasible to find a *collision*: $x \neq y$ such that

$$H(x) = H(y)$$

and infeasible to invert

- add a *nonce* to an email message such that

$$H(\text{nonce} || \text{mail})$$

starts with k zeros

- receiver looks for *nonce* in the message and if not found or incorrect, email is deemed junk

Rationale

- Sender must spend some time before sending out a message
- Receiver must spend (much less) time verifying a message
- Overhead acceptable for the regular email user
- Unfeasible if sending out thousands of email

Cryptographic Hash Functions

- Takes a string x of any length and returns a fixed length output $H(x)$
- Easy to compute (hundreds of megabytes per second)
- Infeasible to find a collision

$$H(x) = H(y)$$

this implies that H is infeasible to invert

- For a given x , only way to find **nonce** such that

$$H(\text{nonce}||x)$$

starts with k zeros is to try different values for **nonce** until one is found. Number of average tries is exponential in k .

Concrete examples: **RipeMD**, **SHA256**, **Keccak** (aka **SHA3**)

Proof of Work in Bitcoin

- Starting from
 - ▶ Hash of previous block
 - ▶ Hash of Transactions of Current Block
 - ▶ Timestamp

Find **nonce** such that the hash value is smaller than **target**.

- Every 2016 blocks (about 2 weeks), the **target** is recomputed by setting **NewTarget** to
$$\text{OldTarget} * \text{MinToProduceLast2016Blocks} / 20160\text{Min}$$
to keep average time between blocks to around 10 minutes

Economic Incentives

A native currency is associated with the Bitcoin blockchain to reward nodes that produce (mine) new blocks


- miner receives the transaction fees associated to each transaction in the block
- a certain number of freshly mint bitcoins
 - ▶ initially 50 bitcoins per block
 - ▶ halving every 210_000 blocks
 - ▶ currently 6.25 bitcoins
 - ▶ **coinbase** transaction, only type of Bitcoin transaction with no input

Extending the blockchain

- At any given time, each node has its own view of the blockchain
- Need not to be one chain
- More chains are possible
- Where to attach the next block?
 - ▶ **Nakamoto consensus**: longest chain available

Forks


Suppose all nodes have the same view of the block chain



Pictures from *Mastering Bitcoin* by Andreas M. Antonopoulos.


Forks

Two blocks found simultaneously




Forks

Two different views of the blockchain coexist




Forks

A new block extends one of fork



Forks

The network converges to one fork and the other is forgotten



Finality

- the transactions in the upside-triangle block disappear from the blockchain and must be resubmitted (unless they are also in the triangle block)
- theoretically the fork can extend to two blocks, if two blocks are found almost simultaneously by nodes on opposite side of the fork
- One-block forks occur about every day
- Two-block forks occur about every few weeks
- It is safe to wait 3 blocks before considering a transaction final
 - ▶ transaction finalization is about 30 min

The Algorand Consensus

Byzantine Agreement

Pease, Shostak and Lamport, 1980

A protocol to reach agreement

Works if at least $2/3$ of the players are honest

- Honey Badger, Miller et al., CCS 2016
 - ▶ A **designated** set of servers reaches consensus through BA
 - ▶ A single point of attack
- Bitcoin-NG, Eyal et al., NSDI 2016
 - ▶ Use Nakamoto consensus to elect a leader
 - ▶ **Leader** publishes transactions
- ByzCoin, Kokoris-Kogias, Usenix Security Symposium, 2016
 - ▶ Use Nakamoto consensus to elect a **group** of participants
- Stellar, D. Mazieres, 2016
 - ▶ Each user **trusts** a subset of other users
 - ▶ Consistent agreement is reached if the transitive closure of the trust relations covers quorum of the nodes

Cryptographic primitives

- Digital Signatures
 - ▶ used to sign messages
- Verifiable Random Functions
 - ▶ used to select users

Digital Signatures

- Digital equivalent of a signature

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document
 - ▶ Infeasible for others to sign on behalf of the legitimate user

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document
 - ▶ Infeasible for others to sign on behalf of the legitimate user
 - ▶ Everybody can verify the signature

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document
 - ▶ Infeasible for others to sign on behalf of the legitimate user
 - ▶ Everybody can verify the signature
- To setup the digital signature, a user executes

$$(\text{sigk}, \text{vk}) \leftarrow \text{GenKey}()$$

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document
 - ▶ Infeasible for others to sign on behalf of the legitimate user
 - ▶ Everybody can verify the signature
- To setup the digital signature, a user executes

$$(\text{sigk}, \text{vk}) \leftarrow \text{GenKey}()$$

- ▶ a **signing key** sigk and a **verification key** vk

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document
 - ▶ Infeasible for others to sign on behalf of the legitimate user
 - ▶ Everybody can verify the signature
- To setup the digital signature, a user executes

$$(\text{sigk}, \text{vk}) \leftarrow \text{GenKey}()$$

- ▶ a **signing key** **sigk** and a **verification key** **vk**
- ▶ the **verification key** is made public (and linked to the identity of the user, if need be)

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document
 - ▶ Infeasible for others to sign on behalf of the legitimate user
 - ▶ Everybody can verify the signature
- To setup the digital signature, a user executes

$$(\text{sigk}, \text{vk}) \leftarrow \text{GenKey}()$$

- ▶ a **signing key** sigk and a **verification key** vk
 - ▶ the **verification key** is made public (and linked to the identity of the user, if need be)
- to sign message m , user executes

$$\text{sig} = \text{Sign}(m, \text{sigk})$$

Digital Signatures

- Digital equivalent of a signature
 - ▶ The legitimate user can sign a document
 - ▶ Infeasible for others to sign on behalf of the legitimate user
 - ▶ Everybody can verify the signature
- To setup the digital signature, a user executes

$$(\text{sigk}, \text{vk}) \leftarrow \text{GenKey}()$$

- ▶ a **signing key** sigk and a **verification key** vk
 - ▶ the **verification key** is made public (and linked to the identity of the user, if need be)
- to sign message m , user executes

$$\text{sig} = \text{Sign}(m, \text{sigk})$$

- given m and sig anybody can check that sig is a signature of m w.r.t. to sigk by executing

$$\text{Verify}(\text{sigk}, \text{sig}, m)$$

Digital Signatures

Digital Signatures in Algorand

- an *address/user* is a *verification key*
- a user *owns* an address, if they know the associated *signing key*
- each Algorand address is associated with a certain number of Algo tokens: the *balance*
- if a user proposes a transaction, i.e. to *transfer/spend* some tokens, the transaction is signed with the *signing key* associated with the address
- everybody can verify that a transaction is properly signed

signing key aka the *spending key*

Algorand uses **Ed25519**: EdDSA with SHA-512 and Curve 25519.

Verifiable Random Functions

[Micali, Rabin, and Vadhan, 1999]

- A Cryptographic Primitive:
 - ▶ a public key pk
 - ▶ a secret key sk
- Given value x and (pk, sk)
 - ▶ it is possible to compute $y = F(sk, x)$ and a proof Π that y is the correct value
 - ▶ for all $x \neq x'$, $y = F(sk, x')$ is still *random* even conditioned on $pk, (x, y)$
- Given (x, y, Π) and a public key pk
 - ▶ it is possible to verify that y is the correct value

▶ VRF used by Algorand \Rightarrow <https://datatracker.ietf.org/doc/html/draft-irtf-cfrg-vrf-10>

The Algorand Consensus

Rationale

- Addresses (i.e., verification keys) are the players
- Each player has a weight proportional to its **stake** (number of Algo tokens associated with the address)
 - ▶ Creating dummy addresses is useless
- Digital Signatures
 - ▶ used to sign messages exchanged during the consensus protocol
- Verifiable Random Functions
 - ▶ used to select players with probability proportional to its stake

Cryptographic Sortition and Participant Replacement

- Cryptographic Sortition

- ▶ Each party evaluates the VRF on some public information of the blockchain to check if selected for a committee
 - ★ if not selected, no action
 - ★ if selected, announce and every one can verify correctness

- Participant Replacement

- ▶ once selected, a participant sends one message
- ▶ no time for an adversary to attack a committee member

Cryptographic Sortition

- τ (expected) number of users to be chosen
- user i with $(\text{pk}_i, \text{sk}_i)$ is selected with probability proportionally to the weight w_i
 - ▶ more than one sub-user $(i, 1), \dots, (i, w_i)$ can be chosen
- for role **role**

Algorithm for Cryptographic Sortition

let s be the **seed** and W total weight

- set $(h, \Pi) = F(\text{sk}_i, s || \text{role})$ and $p = \tau / W$
- return (h, Π, j) such that $h/2^\ell$ falls in interval

$$\left[\sum_{k=0}^j B(k; w_i, p), \quad \sum_{k=0}^{j+1} B(k; w_i, p) \right]$$

where $B(k; w, p) = \binom{w}{k} p^k (1-p)^{w-k}$ is the probability that k out of w_i are selected.

Sybil attacks

Splitting the weight w of a user into $w = w_1 + w_2$ of two dummy users does not help

$$B(k_1; w_1, p) + B(k_2; w_2, p) = B(k_1 + k_2; w_1 + w_2, p)$$

Selecting the seed

- at round $r - 1$, every block proposer u adds also the proposed seed s_r for round r computed as

$$(s, \Pi) = F(\text{sk}_u, s_{r-1-(r \bmod R)} || r)$$

- s_0 chosen at start
- key sk_u must be chosen in advance

Participation Keys

- To take part in consensus, a user must be online
- This means that some secret key of the user should be given to a node
- Users should not use the *spending key* for consensus.
If the node is compromised, the user could lose the Algo tokens
- Instead...

Participation Keys

- a user generates and registers a *participation key* for a certain number of rounds
- it also generates a collection of *ephemeral* keys, one for each round, signs these keys with the participation key, and then deletes the participation key.
- each ephemeral key is used to sign consensus messages for the corresponding round, and is deleted after the round is over.
 - ▶ Using participation keys ensures that a user's algos are secure even if their participating node is compromised.
 - ▶ Deleting the participation and ephemeral keys after they are used ensures that the blockchain is forward-secure and cannot be compromised by attacks on old blocks using old keys.

Consensus

- 1 Block Proposal
- 2 Soft Vote
- 3 Certify Vote.

All messages in each phase are signed with one of the ephemeral keys of the round.

I - Block proposal

- A user runs with cryptographic sortition with $\tau = 26$.
 - ▶ with probability $1 - 10^{-11}$ we have at least one proposer and at most 70 proposers
- Each user propagates block k from $1, \dots, j$ with highest priority

$$H(VRF || k)$$

along with proof and proposed block

I - Block Proposal

- Each node get proposals from other nodes
- Verify the signature of the block and the cryptographic sortition
- Do not propagate blocks with lower priority than already seen
- This is done for a fixed amount of time
 - ▶ tradeoff between efficiency and probability of not receiving any proposed block

II - Soft Vote

- Each node runs cryptographic sortition with `role='committee|round|step'` to see if chosen for the **soft vote committee**
- All chosen accounts will have a weighted vote based on the number of algos the account has
- These votes will be for the highest priority block proposed and will be sent out to the other nodes signed and along with the sortition proof
- If within **timeout** no quorum is formed
 - ▶ A new committee is formed until a **quorum** is reached
 - ▶ for a maximum number of steps

References

- Y. Gilad et al., *Algorand: Scaling Byzantine Agreement for Cryptocurrencies*, SOSP '17
- J. Chen, S. Micali, *Algorand: A secure and efficient distributed ledger*, TCS, 2019
- Algorand Blockchain Feature Specification, 2019
- Algorand Source Code at github.com/algorand/go-algorand

PoW Consensus vs PoS Consensus

PoW vs PoS

- Computational Power

- ▶ **PoW:** very expensive (inverting hash function by brute force)
- ▶ **PoW:** need to be compensated
- ▶ **PoS:** inexpensive (can run on Raspberry pi)
- ▶ **PoS:** no need to be compensated
- ▶ **PoS:** more tokens, more likely to be selected, more work
- ▶ **PoS:** more tokens, greater interest in the network, more willing to be involved

- Robustness to attacks

- ▶ PoW and PoS rely on majority of good players
- ▶ Gain majority to disrupt the networks
- ▶ **PoW:** Buy hardware
- ▶ **PoS:** Buy tokens
 - ★ if you own more than half of the PoS economy
why would you want to disrupt it?